

HW4

February 28, 2019

```
In [1]: import pandas as pd
        from sklearn.linear_model import LinearRegression
        import numpy as np
        import statsmodels.api as sm
```

1 Question #1 (a)

```
In [2]: raw_data = pd.read_csv('Koop-Tobias.csv')
        raw_data.head()
```

```
Out[2]:    PERSONID  EDUC  LOGWAGE  POTEXPER  TIMETRND  ABILITY  MOTHERED  FATHERED \
0           1     13      1.82         1          0       1.0       12        12
1           1     18      3.29         3          7       1.0       12        12
2           1     18      3.21         5          9       1.0       12        12
3           1     18      3.06         6         10       1.0       12        12
4           2     15      2.14         4          6       1.5       12        12

      BRKNHOME  SIBLINGS
0            0        1
1            0        1
2            0        1
3            0        1
4            0        1
```

```
In [3]: X = raw_data.drop(columns=['PERSONID', 'LOGWAGE', 'TIMETRND'])
        X = sm.add_constant(X)
        y = raw_data['LOGWAGE']

        results = sm.OLS(y, X).fit()

        print(results.summary())
```

```
OLS Regression Results
=====
Dep. Variable: LOGWAGE    R-squared:      0.176
Model:          OLS     Adj. R-squared:  0.176
Method:        Least Squares   F-statistic:  546.6
```

```

Date: Wed, 27 Feb 2019   Prob (F-statistic):          0.00
Time: 23:24:58   Log-Likelihood:           -12255.
No. Observations: 17919   AIC:            2.453e+04
Df Residuals:    17911   BIC:            2.459e+04
Df Model:             7
Covariance Type: nonrobust
=====

              coef      std err       t     P>|t|      [0.025      0.975]
-----
const      0.9897     0.034     29.198     0.000      0.923      1.056
EDUC       0.0712     0.002     31.538     0.000      0.067      0.076
POTEXPER    0.0395     0.001     43.970     0.000      0.038      0.041
ABILITY      0.0774     0.005     15.682     0.000      0.068      0.087
MOTHERED    7.099e-05   0.002      0.042     0.967     -0.003      0.003
FATHERED    0.0053     0.001      3.974     0.000      0.003      0.008
BRKNHOME   -0.0529     0.010     -5.292     0.000     -0.072     -0.033
SIBLINGS     0.0049     0.002      2.720     0.007      0.001      0.008
=====
Omnibus:        1110.722   Durbin-Watson:          0.804
Prob(Omnibus):    0.000   Jarque-Bera (JB):      2075.146
Skew:          -0.458   Prob(JB):                 0.00
Kurtosis:        4.393   Cond. No.                  218.
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

marginal effect

```

In [5]: X_bar = X.mean()
         X_bar['BRKNHOME'] = 0
         X_bar['EDUC'] = 12

In [6]: X_bar = np.asarray(X_bar)
         coef = np.asarray(results.params)

In [7]: np.exp(np.matmul(X_bar, coef))*results.params[1]

Out[7]: 0.6800472008090137

```

2 Question #1 (b)

```

In [8]: X['EDUC_2'] = np.power(X['EDUC'], 2)
         X['EDUC_3'] = np.power(X['EDUC'], 3)

         results_2 = sm.OLS(y, X).fit()
         R = np.array(([0,0,0,0,0,0,0,0,1], [0,0,0,0,0,0,0,0,1,0]))

```

```
In [9]: print(results_2.summary())
```

```
OLS Regression Results
=====
Dep. Variable: LOGWAGE R-squared: 0.177
Model: OLS Adj. R-squared: 0.177
Method: Least Squares F-statistic: 428.9
Date: Wed, 27 Feb 2019 Prob (F-statistic): 0.00
Time: 23:24:58 Log-Likelihood: -12240.
No. Observations: 17919 AIC: 2.450e+04
Df Residuals: 17909 BIC: 2.458e+04
Df Model: 9
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	1.3650	0.440	3.102	0.002	0.502	2.228
EDUC	-0.0601	0.099	-0.606	0.545	-0.254	0.134
POTEXPER	0.0396	0.001	44.060	0.000	0.038	0.041
ABILITY	0.0741	0.005	14.908	0.000	0.064	0.084
MOTHERED	0.0003	0.002	0.182	0.855	-0.003	0.004
FATHERED	0.0052	0.001	3.865	0.000	0.003	0.008
BRKNHOME	-0.0499	0.010	-4.990	0.000	-0.070	-0.030
SIBLINGS	0.0048	0.002	2.707	0.007	0.001	0.008
EDUC_2	0.0131	0.007	1.780	0.075	-0.001	0.028
EDUC_3	-0.0004	0.000	-2.214	0.027	-0.001	-4.55e-05
Omnibus:	1114.698	Durbin-Watson:			0.805	
Prob(Omnibus):	0.000	Jarque-Bera (JB):			2107.537	
Skew:	-0.456	Prob(JB):			0.00	
Kurtosis:	4.411	Cond. No.			3.08e+05	

<=====

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.08e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [10]: print(results_2.f_test(R))
```

```
<F test: F=array([[14.34648984]]), p=5.948185372443663e-07, df_denom=17909, df_num=2>
```

3 Question #1 (c)

```
In [11]: X = raw_data['EDUC']
X = sm.add_constant(X)
```

```

y = raw_data['LOGWAGE']
results_3 = sm.OLS(y, X).fit()

y_pred = results_3.predict(X)
lin_params = results_3.params
print(results_3.summary())

OLS Regression Results
=====
Dep. Variable: LOGWAGE R-squared: 0.077
Model: OLS Adj. R-squared: 0.077
Method: Least Squares F-statistic: 1497.
Date: Wed, 27 Feb 2019 Prob (F-statistic): 1.39e-314
Time: 23:24:58 Log-Likelihood: -13270.
No. Observations: 17919 AIC: 2.654e+04
Df Residuals: 17917 BIC: 2.656e+04
Df Model: 1
Covariance Type: nonrobust
=====
            coef    std err          t      P>|t|      [0.025      0.975]
-----
const      1.3296    0.025     52.589      0.000     1.280      1.379
EDUC       0.0763    0.002     38.690      0.000      0.072      0.080
=====
Omnibus:      578.200   Durbin-Watson: 0.793
Prob(Omnibus): 0.000   Jarque-Bera (JB): 878.367
Skew:        -0.319   Prob(JB): 1.84e-191
Kurtosis:      3.878   Cond. No. 86.0
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

```

In [12]: from statsmodels.nonparametric.kernel_regression import KernelReg
X = raw_data['EDUC']
bw = [(np.std(X, ddof=1)* (len(X))**(-1/5))]

results_4 = KernelReg(y, X, 'c', reg_type='lc', bw=bw).fit()
y_pred_np = results_4[0]

```

```

In [15]: import matplotlib.pyplot as plt

fig, ax = plt.subplots()
fig.set_size_inches(18.5, 10.5)

new_x_1, new_y_1 = zip(*sorted(zip(X, y_pred)))

```

```

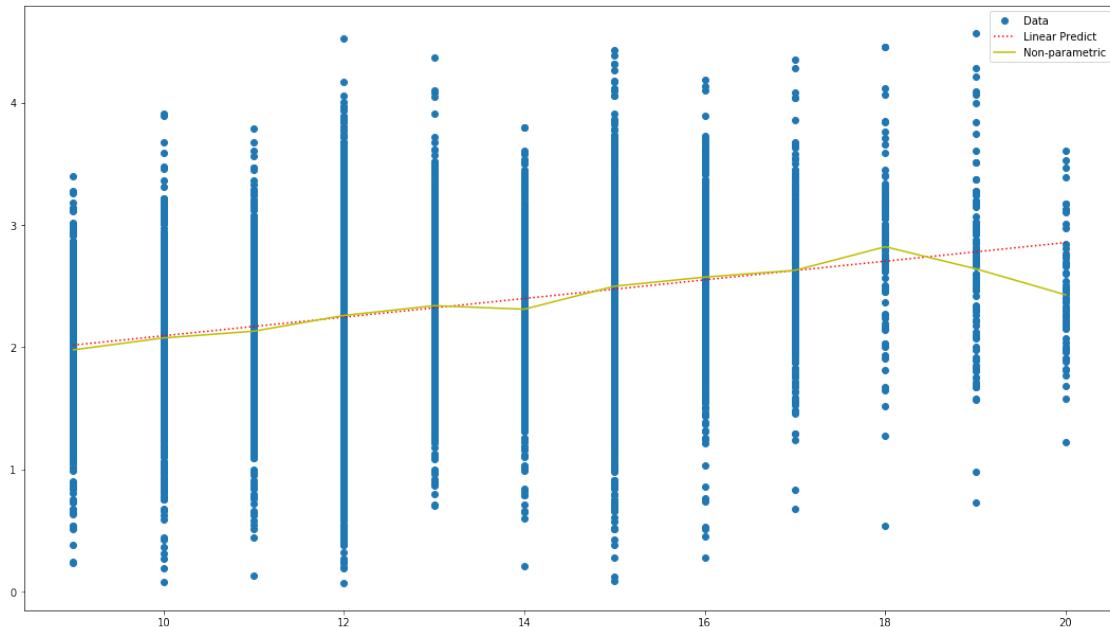
new_x_2, new_y_2 = zip(*sorted(zip(X, y_pred_np)))

ax.plot(X, y, 'o', label="Data")
ax.plot(new_x_1, new_y_1, 'r', label="Linear Predict", linestyle='dotted')
ax.plot(new_x_2, new_y_2, 'y', label="Non-parametric")

ax.legend(loc="best")

```

Out[15]: <matplotlib.legend.Legend at 0x20326690748>



In [16]: marginal_effect = results_4[1]*100

```

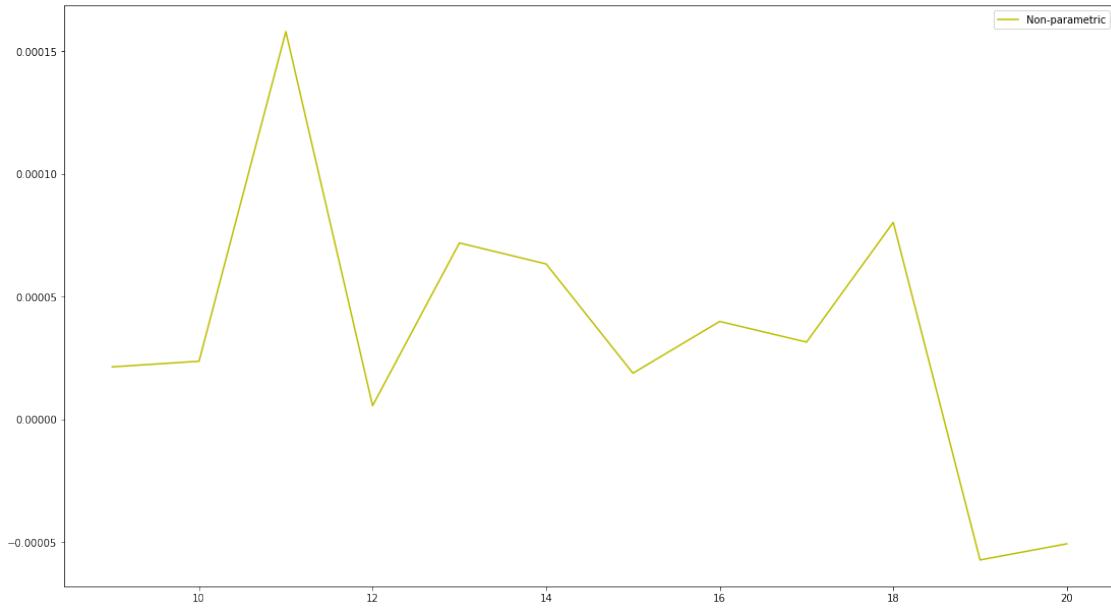
In [18]: fig, ax_2 = plt.subplots()
fig.set_size_inches(18.5, 10.5)
#ax.plot(X, y, 'o', label="Data")
new_x, new_y = zip(*sorted(zip(X, marginal_effect)))

ax_2.plot(new_x, new_y, 'y', label="Non-parametric")
#ax_2.plot(X, np.full(X.shape, lin_params[1]), 'g', label="Linear", linestyle='dashed')

ax_2.legend(loc="best")

```

Out[18]: <matplotlib.legend.Legend at 0x20326717c18>



4 Question #2

```
In [19]: raw_data = pd.read_csv('cps71.csv')

In [20]: X = raw_data['age']
        X = sm.add_constant(X)
        y = raw_data['logwage']

        results = sm.OLS(y, X).fit()
        lin_params = results.params

        y_pred_lin = results.predict(X)

        X['age_2'] = np.power(raw_data['age'], 2)

        results_2 = sm.OLS(y, X).fit()
        quad_params = results_2.params

        y_pred_quad = results_2.predict(X)

In [21]: bw = [(np.std(X['age']), ddof=1)* (len(X))**(-1/5))]
        results_3 = KernelReg(y, X['age'], 'c', reg_type='lc', bw=bw).fit()
        y_pred_np = results_3[0]

In [22]: import matplotlib.pyplot as plt

        fig, ax = plt.subplots()
```

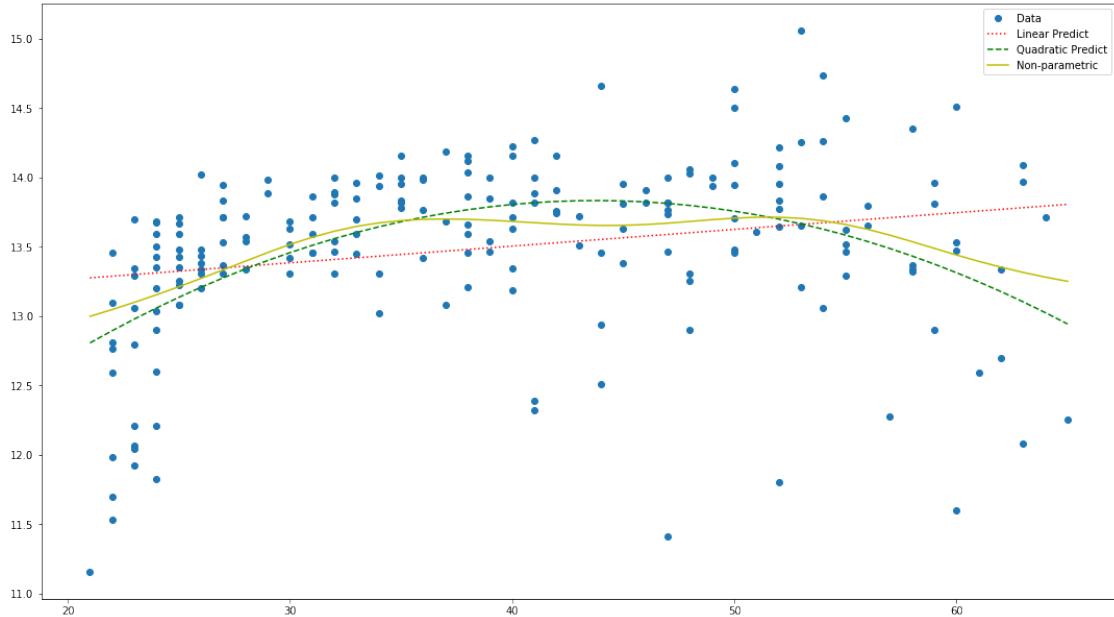
```

fig.set_size_inches(18.5, 10.5)
ax.plot(X['age'], y, 'o', label="Data")
ax.plot(X['age'], y_pred_lin, 'r', label="Linear Predict", linestyle='dotted')
ax.plot(X['age'], y_pred_quad, 'g', label="Quadratic Predict", linestyle='dashed')
ax.plot(X['age'], y_pred_np, 'y', label="Non-parametric")

ax.legend(loc="best")

```

Out [22]: <matplotlib.legend.Legend at 0x20326706a58>



In [23]: np_marginal = results_3[1]

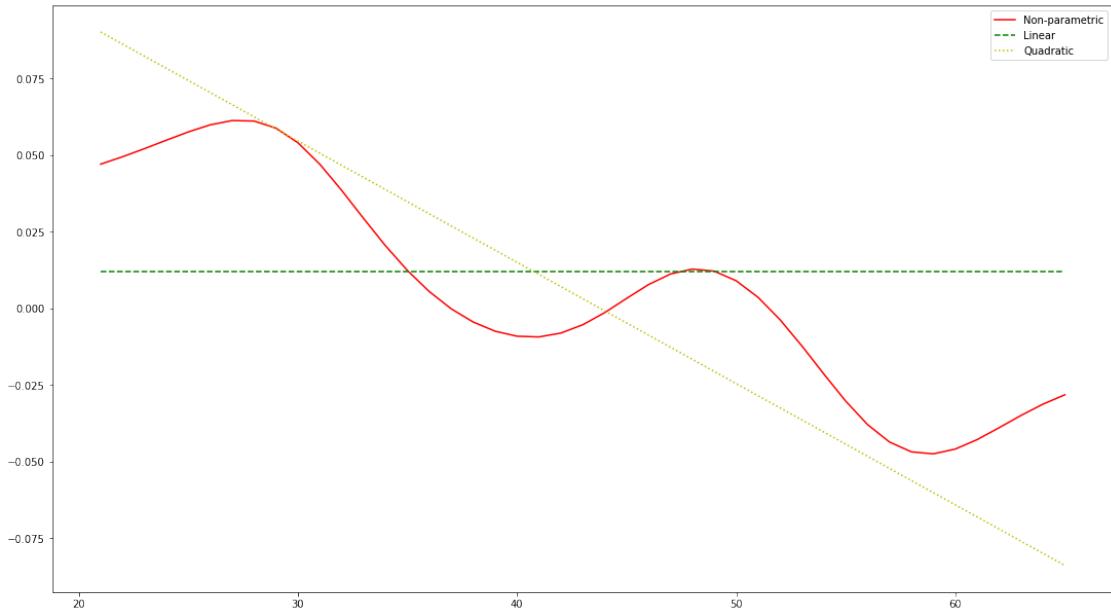
```

fig, ax_2 = plt.subplots()
fig.set_size_inches(18.5, 10.5)
ax_2.plot(X['age'], np_marginal*100, 'r', label="Non-parametric")
ax_2.plot(X['age'], np.full(X['age'].shape, lin_params[1]), 'g', label="Linear", line)
ax_2.plot(X['age'], X['age']*quad_params[2]*2 + quad_params[1], 'y', label="Quadratic")

ax_2.legend(loc="best")

```

Out [23]: <matplotlib.legend.Legend at 0x2032708be80>



5 Question #3(a)

```
In [141]: import math
        sample_sizes = [25, 50, 100] #size of sample
        # y = 4 + 2 * x + error
        # N error N(0,1)
        # T error T(4)
        coef_big_mat = []
        z_big_mat = []
        t_big_mat = []
        for n in sample_sizes:
            coef_n = []
            coef_t = []
            z_n = []
            z_t = []
            t_n = []
            t_t = []
            for t in range(1000):
                n_error = np.random.normal(0, 1, n)
                t_error = np.random.standard_t(4, n)
                x = np.random.uniform(0, 3, n)
                y_n = 4 + 2*x + n_error
                y_t = 4 + 2*x + t_error
                sd = math.sqrt(1/np.sum(np.power(x, 2)))

                # N Error
```

```

        result_n = sm.OLS(y_n, sm.add_constant(x)).fit()
        coef_n.append(result_n.params[1])
        z = (result_n.params[1]-0)/sd
        se = result_n.bse[1]
        t = (result_n.params[1]-0)/se
        t_n.append(t)
        z_n.append(z)

    # T Error
    result_t = sm.OLS(y_t, sm.add_constant(x)).fit()
    coef_t.append(result_t.params[1])
    z = (result_t.params[1]-0)/sd
    se = result_t.bse[1]
    t = (result_t.params[1]-0)/se
    t_t.append(t)
    z_t.append(z)

    coef_big_mat.append([coef_n, coef_t])
    z_big_mat.append([z_n, z_t])
    t_big_mat.append([t_n, t_t])

```

```
In [142]: from matplotlib import pyplot
         from matplotlib.pyplot import figure
```

```
In [143]: bins = np.linspace(-0, 4, 100)
```

```

print("N Error Case:")

figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(coef_big_mat[i][0]))+", var="+str(np.var(coef_big_mat[i][0])))
    pyplot.hist(coef_big_mat[i][0], bins, alpha=0.5, label=str(sample_sizes[i]))
    pyplot.legend(loc='upper right')
    pyplot.show()

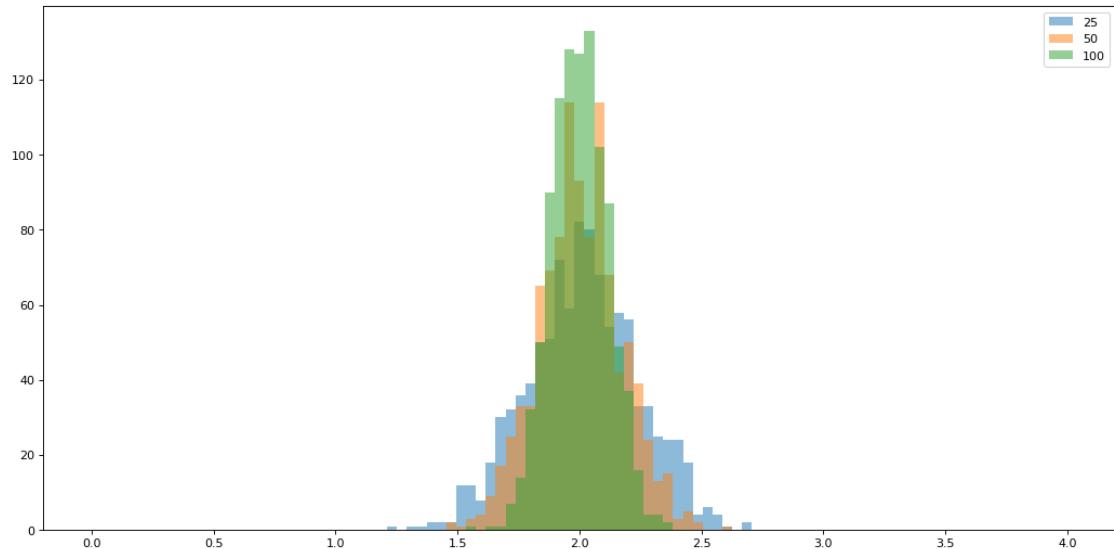
print("T Error Case:")

figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(coef_big_mat[i][1]))+", var="+str(np.var(coef_big_mat[i][1])))
    pyplot.hist(coef_big_mat[i][1], bins, alpha=0.5, label=str(sample_sizes[i]))
    pyplot.legend(loc='upper right')
    pyplot.show()

```

N Error Case:

```
n=25, mean=2.012488983333248, var=0.05249862604346628, se=0.2291257865092148
n=50, mean=2.004038638956284, var=0.028643210231471507, se=0.16924305076271673
n=100, mean=1.9998946283499164, var=0.01354834677651397, se=0.1163973658486908
```

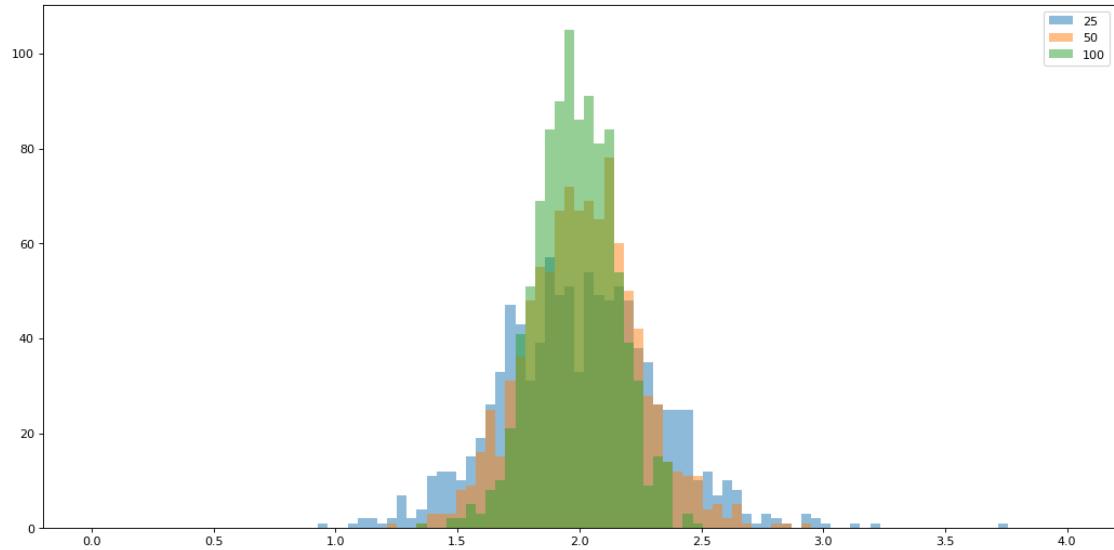


T Error Case:

n=25, mean=2.0060822548285806, var=0.10567402834654102, se=0.32507541947452906

n=50, mean=2.0113618504889645, var=0.05352231792474029, se=0.23134890949546377

n=100, mean=1.9884132856956598, var=0.02630018530277698, se=0.16217331871419843



6 Question #3(b)

6.0.1 Z Statistics

```
In [144]: bins = np.linspace(5, 55, 100)
```

```
print("N Error Case:")

figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(z_big_mat[i][0]))+", var="+str(np.var(z_big_mat[i][0])))
    count = 0
    for element in z_big_mat[i][0]:
        if (element > 1.96 or element < -1.96):
            count += 1
    print("size of test:", 1 - count/1000)
    pyplot.hist(z_big_mat[i][0], bins, alpha=0.5, label=str(sample_sizes[i]))
    pyplot.legend(loc='upper right')
    pyplot.show()

print("T Error Case:")

figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(z_big_mat[i][1]))+", var="+str(np.var(z_big_mat[i][1])))
    count = 0
    for element in z_big_mat[i][1]:
        if (element > 1.96 or element < -1.96):
            count += 1
    print("size of test:", 1 - count/1000)
    pyplot.hist(z_big_mat[i][1], bins, alpha=0.5, label=str(sample_sizes[i]))
    pyplot.legend(loc='upper right')
    pyplot.show()
```

N Error Case:

n=25, mean=17.268458433839506, var=6.4753934834130495

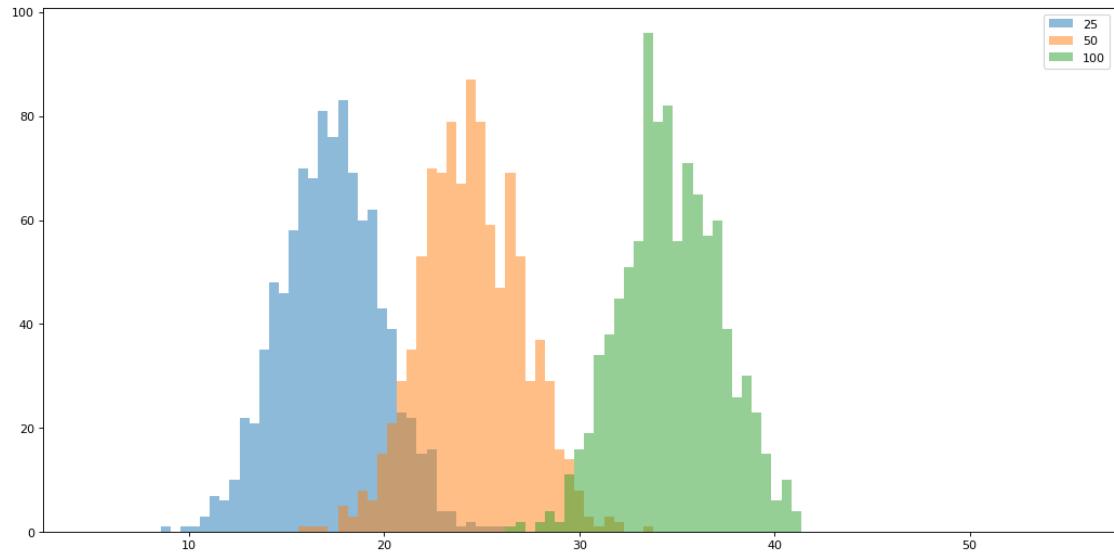
size of test: 0.0

n=50, mean=24.478947133064747, var=6.603008365729441

size of test: 0.0

n=100, mean=34.689623939681276, var=6.482227300655119

size of test: 0.0



T Error Case:

n=25, mean=17.21294748809783, var=10.290138070588648

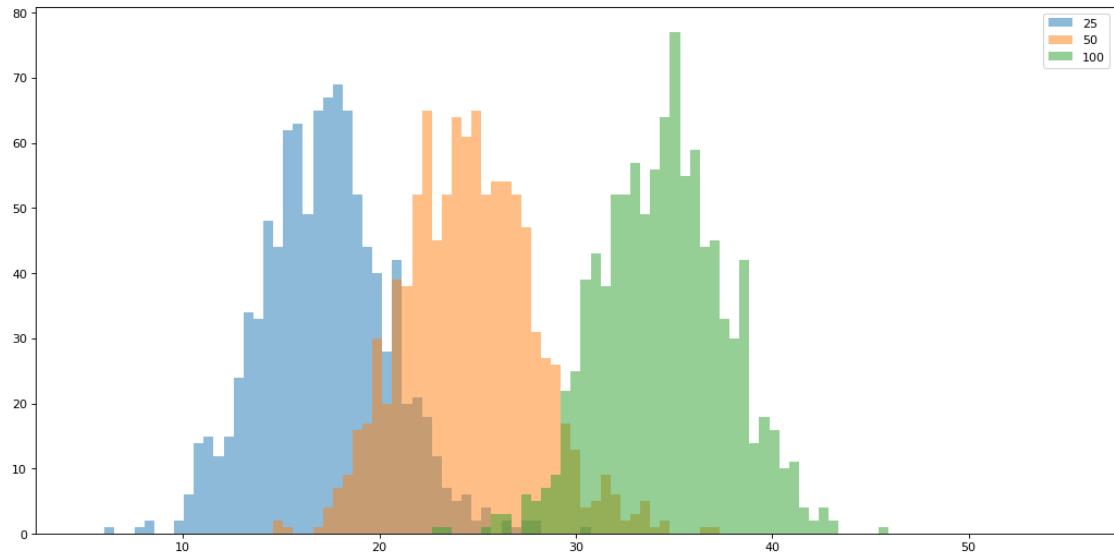
size of test: 0.0

n=50, mean=24.57051460244579, var=10.509378822227127

size of test: 0.0

n=100, mean=34.48750148998313, var=10.095753493509088

size of test: 0.0



6.0.2 T Statistics

```
In [145]: bins = np.linspace(5, 45, 100)

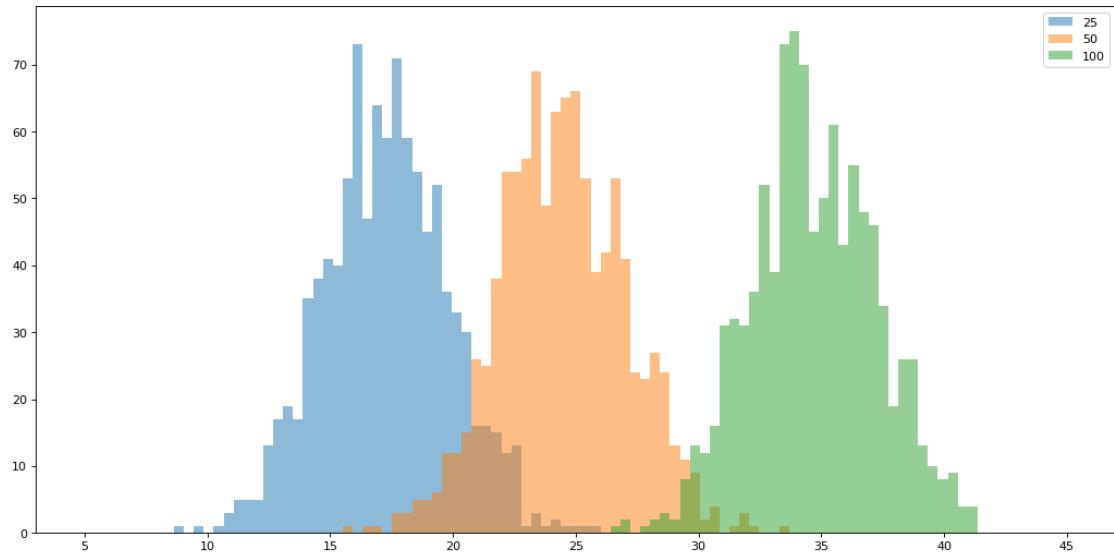
print("N Error Case:")

figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(t_big_mat[i][0]))+", var="
    count = 0
    for element in t_big_mat[i][0]:
        if (element > 12.706 or element < -12.706):
            count += 1
    print("size of test:", 1 - count/1000)
    pyplot.hist(z_big_mat[i][0], bins, alpha=0.5, label=str(sample_sizes[i]))
pyplot.legend(loc='upper right')
pyplot.show()

print("T Error Case:")

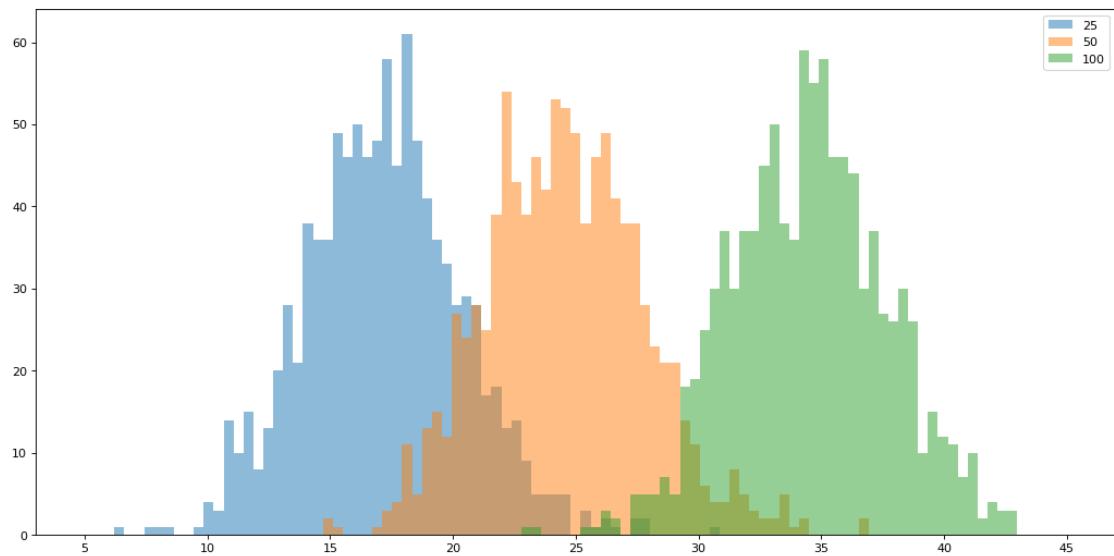
figure(num=None, figsize=(16, 8), dpi=80, facecolor='w', edgecolor='k')
for i in range(3):
    print("n="+str(sample_sizes[i])+", mean="+str(np.mean(t_big_mat[i][1]))+", var="
    count = 0
    for element in t_big_mat[i][1]:
        if (element > 12.706 or element < -12.706):
            count += 1
    print("size of test:", 1 - count/1000)
    pyplot.hist(z_big_mat[i][1], bins, alpha=0.5, label=str(sample_sizes[i]))
pyplot.legend(loc='upper right')
pyplot.show()

N Error Case:
n=25, mean=8.735152019547114, var=3.3629344001402646
size of test: 0.968
n=50, mean=12.228941819914041, var=3.263578798541486
size of test: 0.635
n=100, mean=17.351736876535266, var=3.5030257513385066
size of test: 0.00200000000000000018
```



T Error Case:

```
n=25, mean=6.627049466024269, var=4.000635431527382
size of test: 0.991
n=50, mean=9.186082838319468, var=4.161971560960327
size of test: 0.957
n=100, mean=12.504484595267229, var=4.389229381841637
size of test: 0.5329999999999999
```



7 Question #3(c)

```
In [146]: sample_sizes = [25, 50, 100] #size of sample
    #  $y = 4 + 2 * x + error$ 
    print("N Error Case:")
    sample_X = np.random.uniform(0, 3, 25)
    sample_n_error = np.random.normal(0, 1, 25)
    sample_y_n = 4 + 2*sample_X + sample_n_error
    result_original = sm.OLS(sample_y_n, sm.add_constant(sample_X)).fit()
    beta_1 = result_original.params[0]
    beta_2 = result_original.params[1]
    u_hat = sample_y_n - beta_1 - beta_2 * sample_X

    coef = []
    for t in range(1000):
        u_star = np.random.choice(u_hat, 25)
        y_star = beta_1 + beta_2 * sample_X + u_star
        result_bootstrap = sm.OLS(y_star, sm.add_constant(sample_X)).fit()
        coef.append(result_bootstrap.params[1])
    print("sample beta:", str(beta_2))
    print("bootstrap beta:", str(np.mean(coef)))
    print("bootstrap se:", str(math.sqrt(np.var(coef))))


print("T Error Case:")
sample_X = np.random.uniform(0, 3, 25)
sample_t_error = np.random.standard_t(4, 25)
sample_y_t = 4 + 2*sample_X + sample_t_error
result_original = sm.OLS(sample_y_t, sm.add_constant(sample_X)).fit()
beta_1 = result_original.params[0]
beta_2 = result_original.params[1]
u_hat = sample_y_n - beta_1 - beta_2 * sample_X

coef = []
for t in range(1000):
    u_star = np.random.choice(u_hat, 25)
    y_star = beta_1 + beta_2 * sample_X + u_star
    result_bootstrap = sm.OLS(y_star, sm.add_constant(sample_X)).fit()
    coef.append(result_bootstrap.params[1])
print("sample beta:", str(beta_2))
print("bootstrap beta:", str(np.mean(coef)))
print("bootstrap se:", str(math.sqrt(np.var(coef))))
```

N Error Case:

```
sample beta: 1.7102102215594508
bootstrap beta: 1.7159064524085028
bootstrap se: 0.1928669790849424
```

T Error Case:

```
sample beta: 1.2898398249410796  
bootstrap beta: 1.2824862775758408  
bootstrap se: 0.5631121845139818
```